

spark-Streaming 读取kafka 在zookeeper中手动维护offset

实验说明:

在kafka集群中创建以offsettest主题，启动生产者进程生产数据，利用在zookeeper维护一个以offsetTest命名的消费者组的offset偏移量

1. offset 在zookeeper中存储位置

首先启动zookeeper客户端

```
atguigu@hadoop101 zookeeper-3.4.10]$ bin/zkCli.sh
```

```
WatchedEvent state:SyncConnected type:None path:null
[zk: localhost:2181(CONNECTED) 0] ls /consumers
[console-consumer-41980, test, console-consumer-23445, offsetTest]
[zk: localhost:2181(CONNECTED) 1] ls /consumers/offsetTest
[offsets]
[zk: localhost:2181(CONNECTED) 2] ls /consumers/offsetTest
[offsets]
[zk: localhost:2181(CONNECTED) 3] ls /consumers/offsetTest/offsets
[offsettest]
[zk: localhost:2181(CONNECTED) 4] ls /consumers/offsetTest/offsets/offsettest
[1, 0]
[zk: localhost:2181(CONNECTED) 5] get /consumers/offsetTest/offsets/offsettest/0
00t18
cZxid = 0x19000000090
ctime = Mon Oct 21 19:31:04 CST 2019
mZxid = 0x1a00000105
mtime = Tue Oct 22 16:52:24 CST 2019
pZxid = 0x19000000090
version = 0
dataVersion = 59
aclVersion = 0
ephemeralOwner = 0x0
dataLength = 9
numChildren = 0
```

进入客户端后输入命令

```
ls /consumers
```

consumer文件夹中是zookeeper中维护的消费者组。

注：当然我们第一次运行程序的时候，在zookeeper中是没有offsetTest消费者组的信息的。新版本已经不维护在zookeeper中，维护在broker集群中的__consumer_offset中，所以在zookeeper中没有消费者组的信息

假设我们已经运行了一此程序，在zookeoer成功创建了offsetTest消费者组

在offsetTest消费者组维护的是offset，offset文件夹下则维护的是各个主题的文件夹，每个主题的文件夹下又是各个分区的信息，在每个分区的信息中第一行就是维护的偏移量。各个分区的信息如下

cZxid: 节点创建时的zxid

ctime: 节点创建时间

mZxid: 节点最近一次更新时的zxid

mtime: 节点最近一次更新的时间

cversion: 子节点数据更新次数

dataVersion: 本节点数据更新次数

aclVersion: 节点ACL(授权信息)的更新次数

ephemeralOwner: 如果该节点为临时节点,ephemeralOwner值表示与该节点绑定的session id. 如果该节点不是临时节点,ephemeralOwner值为0

dataLength: 节点数据长度,本例中为hello world的长度

numChildren: 子节点个数

因此得到在zookeeper中最终的offset偏移量地址

```
/consumers/offsetTest/offsets/offsettest/各个分区: [0,1,2,...]
```

```
/消费者/各个消费者组文件夹/偏移量文件夹/各个topic文件夹
```

2. 实战代码

```
package com.bupt.sparkStreaming

import kafka.common.TopicAndPartition
import kafka.message.MessageAndMetadata
import kafka.serializer.StringDecoder
import kafka.utils.{ZKGroupTopicDirs, ZkUtils}
import org.I0Itec.zkclient.ZkClient
import org.apache.kafka.clients.consumer.ConsumerConfig
import org.apache.log4j.{Level, Logger}
import org.apache.spark.SparkConf
import org.apache.spark.streaming.{Seconds, StreamingContext}
import org.apache.spark.streaming.dstream.InputDStream
import org.apache.spark.streaming.kafka.{HasOffsetRanges, KafkaUtils}

object KafkaDirect_ZK_Offset {
  def main(args: Array[String]): Unit = {
    Logger.getLogger("org.apache.spark").setLevel(Level.OFF)
    val conf: SparkConf = new
    SparkConf().setAppName("KafkaDirect_ZK_Offset").setMaster("local[*]")
    val ssc: StreamingContext = new StreamingContext(conf, Seconds(5))
    val groupId = "offsetTest"
    val brokers = "hadoop101:9092"
    /**
     * kafka参数列表
     */
    val kafkaParams = Map[String, String](
      ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG -> brokers,
      ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG ->
      "org.apache.kafka.common.serialization.StringDeserializer",
      ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG ->
      "org.apache.kafka.common.serialization.StringDeserializer",
```

```

ConsumerConfig.GROUP_ID_CONFIG -> groupId
)

val topic = "offsettest"
// val topics = ArrayList(topic)

val zkGroupTopicDirs: ZKGroupTopicDirs = new ZKGroupTopicDirs(groupId,
topic)
/**
 * 生成的目录结构
 * /customer/offsetTest/offsets/offsettest
 */

val offsetDir: String = zkGroupTopicDirs.consumerOffsetDir
//zk字符串连接组
// val zkGroups = "hadoop101:2181,hadoop102:2181"
val zkGroups = "192.168.1.102:2181"
println(offsetDir)
//创建一个zkClient连接
val zkClient: ZkClient = new ZkClient(zkGroups)
//子节点的数量
val childrenCount: Int = zkClient.countChildren(offsetDir)

//子节点的数量>0就说明非第一次
var stream: InputDStream[(String, String)] = null
if (childrenCount > 0) {
    println("已经启动过")
    //用来存储我们已经读取到的偏移量
    var fromOffset = Map[TopicAndPartition, Long]()
    (0 until childrenCount).foreach(partitionId => {
        val offset = zkClient.readData[String](offsetDir + s"/$partitionId")
        fromOffset += (new TopicAndPartition(topic, partitionId) ->
offset.toLong)
    })
    val mess = (mmd: MessageAndMetadata[String, String])=>
(mmd.key(),mmd.message()) //这个会将 kafka 的消息进行 transform, 最终 kafka的数据都会
变成 (topic_name, message) 这样的 tuple
    stream =
KafkaUtils.createDirectStream[String,String,StringDecoder,StringDecoder,
(String,String)](ssc, kafkaParams,fromOffset,mess)
/*      for (tp <- fromOffset) {
        println(s"topAndPartition: ${tp._1.toString}  offset: ${tp._2}")
      }*/
}
else {
    println("第一次启动")
    stream = KafkaUtils.createDirectStream[String, String, StringDecoder,
StringDecoder](ssc, kafkaParams, Set(topic))

}

val flatMapStream = stream.flatMap(t => t._2.split(" "))
val mapWithStateDStream = flatMapStream.map(_._1)
var result = mapWithStateDStream.reduceByKey(_+_ )
// result.print()
stream.foreachRDD(
    rdd => {
        //转换rdd为Array[OffsetRange]

```

```

val offsetRanges = rdd.asInstanceOf[HasOffsetRanges].offsetRanges

//计算逻辑

//      rdd.foreach(println)
//      rdd.partitions.foreach(println)
//      rdd.values.foreach(println)
//自己存储数据，自己管理
for (o <- offsetRanges) {
    //写入到zookeeper,第二个参数为是否启动安全
    println(s"topicAndPartition: ${o.topicAndPartition()} top: ${o.topic}
patition: ${o.partition} utiloffset: ${o.untilOffset} fromoffset:
${o.fromOffset}")
    val path = offsetDir + "/" + o.partition
    //    println(path)
    zkUtils.updatePersistentPath(zkClient, path, o.untilOffset.toString)
}

)
ssc.start()
ssc.awaitTermination()
}
}

```

pom

```

<dependency>
    <groupId>org.apache.spark</groupId>
    <artifactId>spark-streaming_2.10</artifactId>
    <version>2.1.0</version>
</dependency>

<dependency>
    <groupId>org.apache.spark</groupId>
    <artifactId>spark-streaming-kafka-0-8_2.10</artifactId>
    <version>2.2.0</version>
</dependency>
<!--<dependency>-->
    <!--<groupId>org.apache.spark</groupId>-->
    <!--<artifactId>spark-streaming-kafka_2.10</artifactId>-->
    <!--<version>1.0.0</version>-->
<!--</dependency>-->
<dependency>
    <groupId>org.apache.kafka</groupId>
    <artifactId>kafka-clients</artifactId>
    <version>0.11.0.0</version>
</dependency>

```